

For this problem we use Image Segmentation. An image is a collection or set of different pixels. We group together the pixels that have similar attributes using image segmentation.. Thus, the task of image segmentation is to train a neural network to output a pixel-wise mask of the image. This helps in understanding the image at a much lower level.

for this problem we are going to use Semantic Segmentation U-Net model

https://scikit-image.org/docs/stable/user_guide/tutorial_segmentation.html

```
import numpy as np
import pandas as pd
import tensorflow as tf
from zipfile import ZipFile
import keras.backend as K
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import os

train_zip = "/content/drive/MyDrive/DUTS-TR-Image.zip"
with ZipFile(train_zip, 'r') as zip_:
    zip_.extractall('content')

train_mask_zip = "/content/drive/MyDrive/DUTS-TR-Mask.zip"
with ZipFile(train_mask_zip, 'r') as zip_:
    zip_.extractall('content')

print("Image: ", len(os.listdir("/content/content/DUTS-TR-Image")))
print("Masks:", len(os.listdir("/content/content/DUTS-TR-Mask")))

Image: 10553
Masks: 10553

from IPython.core.display import Image
image = []
paths = []
for dirname, _, filenames in os.walk('/content/content/DUTS-TR-Image'):
    for filename in filenames:
        path = os.path.join(dirname, filename)
        paths.append(path)

        image_id = filename.split(".")[0]
        image.append(image_id)

d = {"id": image, "image_path": paths}
```

```
df = pd.DataFrame(data = d)
df = df.set_index('id')
df.head()
```

image_path



id	
ILSVRC2013_test_00008612	/content/content/DUTS-TR-Image/ILSVRC2013_test...
ILSVRC2014_train_00019927	/content/content/DUTS-TR-Image/ILSVRC2014_trai...
n07714571_1565	/content/content/DUTS-TR-Image/n07714571_1565.jpg
sun_acswpggpfdacoxzay	/content/content/DUTS-TR-Image/sun_acswpggpfdac...
n02701002_4128	/content/content/DUTS-TR-Image/n02701002_4128.jpg

```
from IPython.core.display import Image
image = []
mask_path = []
for dirname, _, filenames in os.walk('/content/content/DUTS-TR-Mask'):
    for filename in filenames:
        path = os.path.join(dirname, filename)
        mask_path.append(path)

        image_id = filename.split(".")[0]
        image_id = image_id.split("-Mask")[0]
        image.append(image_id)

d = {"id": image, "mask_path": mask_path}
mask_df = pd.DataFrame(data = d)
mask_df = mask_df.set_index('id')
mask_df.head()
```

mask_path



id	
n03775546_8657	/content/content/DUTS-TR-Mask/n03775546_8657.png
ILSVRC2012_test_00020037	/content/content/DUTS-TR-Mask/ILSVRC2012_test_...
ILSVRC2013_test_00007910	/content/content/DUTS-TR-Mask/ILSVRC2013_test_...
n07714571_17316	/content/content/DUTS-TR-Mask/n07714571_17316.png
n03710721_2131	/content/content/DUTS-TR-Mask/n03710721_2131.png

```
df["mask_path"] = mask_df["mask_path"]
df.head()
```

	image_path	mask_path
id		
ILSVRC2013_test_00008612	/content/content/DUTS-TR-Image/ILSVRC2013_test...	/content/content/DUTS-TR-Mask/ILSVRC2013_test_...
ILSVRC2014_train_00019927	/content/content/DUTS-TR-Image/ILSVRC2014_train...	/content/content/DUTS-TR-Mask/ILSVRC2014_train...
n07714571_1565	/content/content/DUTS-TR-Image/n07714571_1565.jpg	/content/content/DUTS-TR-Mask/n07714571_1565.png

Now we use normalize the image pixel in between 0 and 1.

```

image_path = "/content/content/DUTS-TR-Image/ILSVRC2013_test_00008612.jpg"
mask_path = "/content/content/DUTS-TR-Mask/ILSVRC2013_test_00008612.png"

img_size = [256,256]

def data_augmentation(img, mask_img):

    if tf.random.uniform(()) > 0.5:
        img = tf.image.flip_left_right(img)
        mask_img = tf.image.flip_left_right(mask_img)

    return img, mask_img

def preprocessing(image_path, mask_path):
    img = tf.io.read_file(image_path)
    img = tf.image.decode_jpeg(img, channels=3)
    img = tf.image.resize(img, img_size)
    img = tf.cast(img, tf.float32) / 255.0

    mask_img = tf.io.read_file(mask_path)
    mask_img = tf.image.decode_jpeg(mask_img, channels=3)
    mask_img = tf.image.resize(mask_img, img_size)
    mask_img = mask_img[:, :, :1]
    mask_img = tf.math.sign(mask_img)

    return img, mask_img

def create_dataset(df, train = False):
    if not train:
        ds = tf.data.Dataset.from_tensor_slices((df["image_path"].values, df["mask_path"].values))
        ds = ds.map(preprocessing, tf.data.AUTOTUNE)
    else:
        ds = tf.data.Dataset.from_tensor_slices((df["image_path"].values, df["mask_path"].values))
        ds = ds.map(preprocessing, tf.data.AUTOTUNE)
        ds = ds.map(data_augmentation, tf.data.AUTOTUNE)

    return ds

df.head()
```

	image_path	mask_path
id		
ILSVRC2013_test_00008612	/content/content/DUTS-TR-Image/ILSVRC2013_test...	/content/content/DUTS-TR-Mask/ILSVRC2013_test_...
ILSVRC2014_train_00019927	/content/content/DUTS-TR-Image/ILSVRC2014_train...	/content/content/DUTS-TR-Mask/ILSVRC2014_train...
n07714571_1565	/content/content/DUTS-TR-Image/n07714571_1565.jpg	/content/content/DUTS-TR-Mask/n07714571_1565.png
sun_acswpgpfdacoxzay	/content/content/DUTS-TR-Image/sun_acswpgpfdacoxzay	/content/content/DUTS-TR-Mask/sun_acswpgpfdacoxzay

#Splitting dataset into train and test

```
train_df, valid_df = train_test_split(df, random_state=42, test_size=.25)
```

```
train = create_dataset(train_df, train = True)
```

```
valid = create_dataset(valid_df)
```

```
print(train_df.shape)
```

```
print(valid_df.shape)
```

```
(7914, 2)
```

```
(2639, 2)
```

```
TRAIN_LENGTH = len(train_df)
```

```
BATCH_SIZE = 16
```

```
BUFFER_SIZE = 1000
```

```
train_dataset = train.cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE).repeat()
```

```
train_dataset = train_dataset.prefetch(buffer_size=tf.data.AUTOTUNE)
```

```
valid_dataset = valid.batch(BATCH_SIZE)
```

```
def display(display_list):
```

```
    plt.figure(figsize=(15, 15))
```

```
    title = ['Input Image', 'True Mask', 'Predicted Mask']
```

```
    for i in range(len(display_list)):
```

```
        plt.subplot(1, len(display_list), i+1)
```

```
        plt.title(title[i])
```

```
        plt.imshow(tf.keras.preprocessing.image.array_to_img(display_list[i]))
```

```
        plt.axis('off')
```

```
    plt.show()
```

```
#show image corresponding to the mask
```

```
for i in range(5):
```

```
    for image, mask in train.take(i):
```


```
        sample_image, sample_mask = image, mask
```

```
        display([sample_image, sample_mask])
```

Input Image




True Mask



Input Image




True Mask



Input Image



True Mask



Input Image



True Mask






Input Image



True Mask



Input Image




True Mask



Input Image



True Mask






Input Image



True Mask



Input Image





True Mask



Input Image



True Mask



Paper: <https://blog.paperspace.com/unet-architecture-image-segmentation/>

Explanation of U-net: <https://www.youtube.com/watch?v=yG6GbEtGUrU>

▼ Model--we are going to use U-Net model

```
base_model = tf.keras.applications.MobileNetV2(input_shape=[256, 256, 3], include_top=False)
```

```
# Use the activations of these layers
```

```
layer_names = [
```

```
    'block_1_expand_relu',    # 64x64
```

```
    'block_3_expand_relu',    # 32x32
```

```
    'block_6_expand_relu',    # 16x16
```

```
    'block_13_expand_relu',   # 8x8
```

```
    'block_16_project',       # 4x4
```

```
]
```


```
base_model_outputs = [base_model.get_layer(name).output for name in layer_names]
```

```
# Create the feature extraction model
```

```
down_stack = tf.keras.Model(inputs=base_model.input, outputs=base_model_outputs)
```

```
down_stack.trainable = False
```

```
WARNING:tensorflow: `input_shape` is undefined or non-square, or `rows` is not in [96
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/m9412608/9406464 [=====] - 0s 0us/step
9420800/9406464 [=====] - 0s 0us/step
```



```
#use residual connection
```

```
def upsample(filters, size, norm_type='batchnorm', apply_dropout=False):
```

```
    initializer = tf.random_normal_initializer(0., 0.02)
```

```
    result = tf.keras.Sequential()
```

```
    result.add(
```

```
        tf.keras.layers.Conv2DTranspose(filters, size, strides=2,
                                         padding='same',
                                         kernel_initializer=initializer,
                                         use_bias=False))
```

```
    if norm_type.lower() == 'batchnorm':
```

```
        result.add(tf.keras.layers.BatchNormalization())
```

```
    elif norm_type.lower() == 'instancenorm':
```

```
        result.add(tf.keras.layers.InstanceNormalization())
```

```
    if apply_dropout:
```

```
        result.add(tf.keras.layers.Dropout(0.5))
```

```
    result.add(tf.keras.layers.ReLU())
```

```

    return result

up_stack = [
    upsample(512, 3), # 4x4 -> 8x8
    upsample(256, 3), # 8x8 -> 16x16
    upsample(128, 3), # 16x16 -> 32x32
    upsample(64, 3), # 32x32 -> 64x64
]

#adding layerouts
def unet_model(output_channels):
    inputs = tf.keras.layers.Input(shape=[256, 256, 3])

    # Downsampling through the model
    skips = down_stack(inputs)
    x = skips[-1]
    skips = reversed(skips[:-1])

    # Upsampling and establishing the skip connections
    for up, skip in zip(up_stack, skips):
        x = up(x)
        concat = tf.keras.layers.Concatenate()
        x = concat([x, skip])

    # This is the last layer of the model
    last = tf.keras.layers.Conv2DTranspose(
        output_channels, 3, strides=2, activation='sigmoid',
        padding='same') #64x64 -> 128x128

    x = last(x)

    return tf.keras.Model(inputs=inputs, outputs=x)

```

▼ Now We are Train the Model

```

def dice_coef(y_true, y_pred, smooth=1):
    intersection = K.sum(y_true * y_pred, axis=[1,2,3])
    union = K.sum(y_true, axis=[1,2,3]) + K.sum(y_pred, axis=[1,2,3])
    return K.mean( (2. * intersection + smooth) / (union + smooth), axis=0)

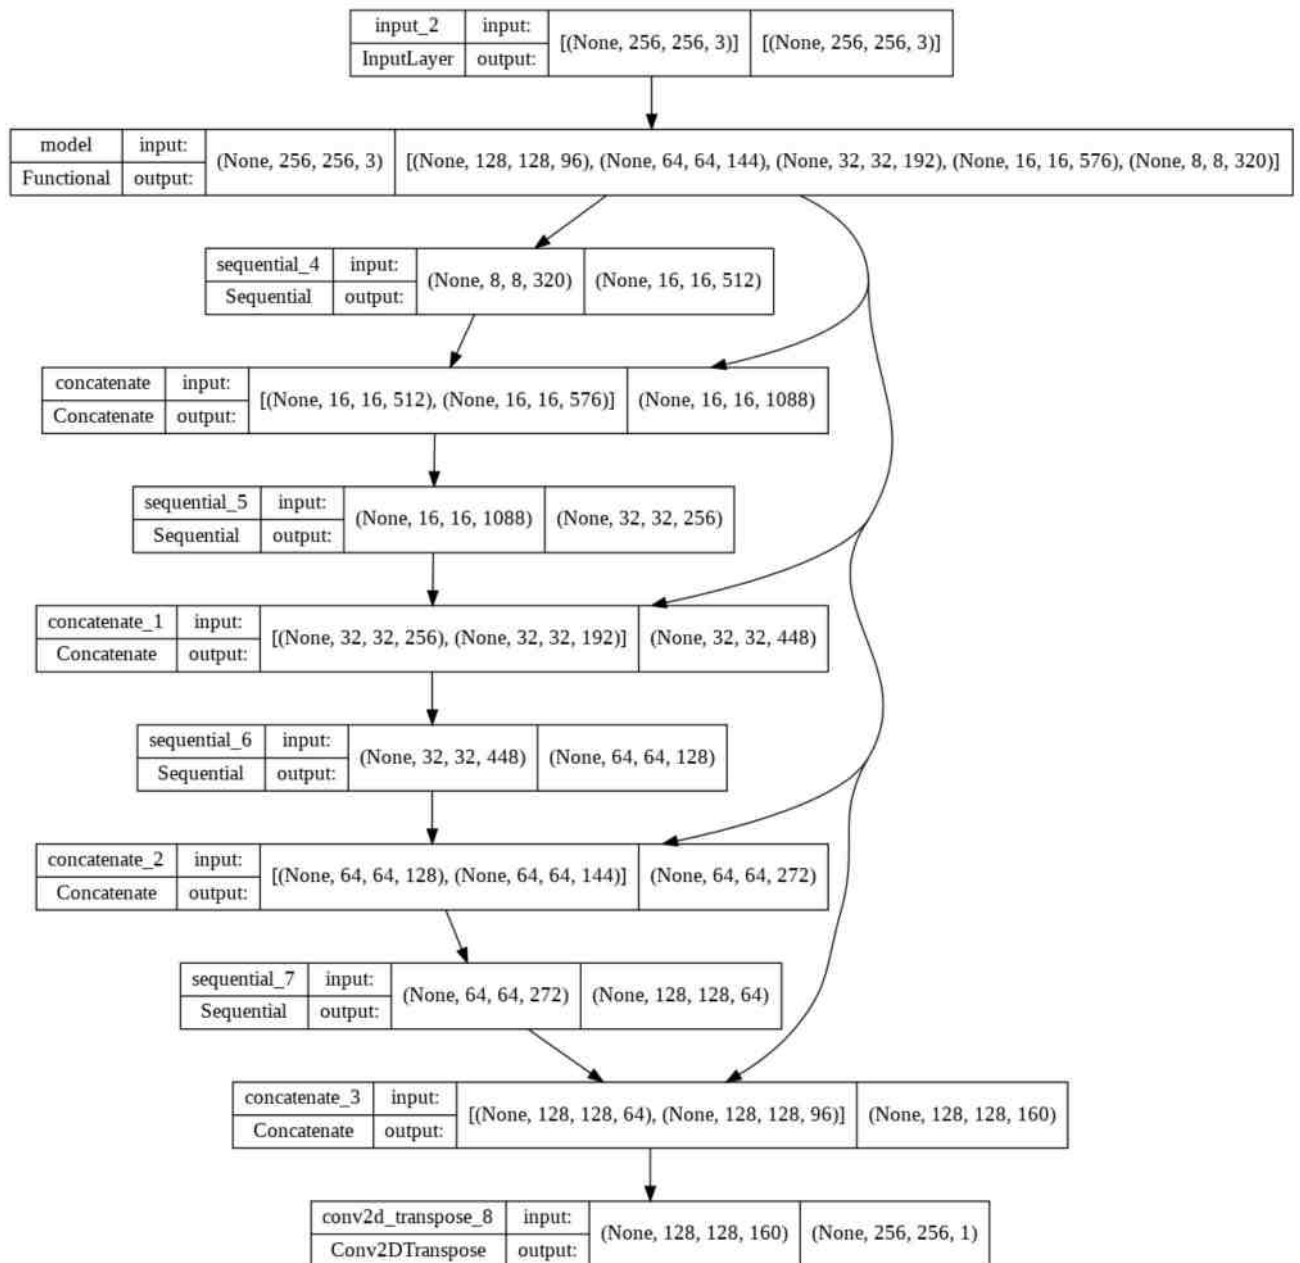
def dice_loss(in_gt, in_pred):
    return 1-dice_coef(in_gt, in_pred)

model = unet_model(1)

model.compile(optimizer='adam',
              loss = dice_loss,
              metrics=[dice_coef, 'binary_accuracy'])

tf.keras.utils.plot_model(model, show_shapes=True)

```



Let's try out the model to see what it predicts before training.

```
for images, masks in train_dataset.take(1):
    for img, mask in zip(images, masks):
```


```

for img, mask in zip(images, masks):
    sample_image = img
    sample_mask = mask
    break
def visualize(display_list):
    plt.figure(figsize=(15, 15))
    title = ['Input Image', 'True Mask', 'Predicted Mask']
    for i in range(len(display_list)):
        plt.subplot(1, len(display_list), i+1)
        plt.title(title[i])
        plt.imshow(tf.keras.preprocessing.image.array_to_img(display_list[i]))
        plt.axis('off')
    plt.show()

def show_predictions(sample_image, sample_mask):
    pred_mask = model.predict(sample_image[tf.newaxis, ...])
    pred_mask = pred_mask.reshape(img_size[0],img_size[1],1)
    visualize([sample_image, sample_mask, pred_mask])

show_predictions(sample_image, sample_mask)

```



```
model.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 256, 256, 3)]	0	[]
model (Functional)	[(None, 128, 128, 9 6), (None, 64, 64, 144) (None, 32, 32, 192) (None, 16, 16, 576) (None, 8, 8, 320)]	1841984	['input_2[0][0]']
sequential_4 (Sequential)	(None, 16, 16, 512)	1476608	['model[0][4]']